

In this paper, (Meng et al., 2022b) define $m_{[t]}^l = W_{\text{out}}^l(\sigma(W_{\text{in}}^l \gamma(h_{[t]}^{l-1})))$. Given this, define

$$m_i^l = W_{\text{out}}^l k_i^l + r_i^l$$

$$\text{where } r_i^l = \frac{z_i - h_i^L}{L - l + 1} \quad (8)$$

Note that the denominator of r_i^l allows us to spread out the burden across multiple layer, allowing for a more scalable algorithm. It is hard to compute C^l exactly, however it can be reliably estimated using $C^l = \lambda \mathbb{E}_k[k^l k^{l\top}]$ over randomly sampled inputs, where $\lambda = 1.5 \times 10^4$. Incorporating the update gives us our desired new weights \hat{W}_{out}^l

A.6.3 MEND

Using the fact that the gradient of loss L with respect to the weights W_ℓ of layer ℓ of an MLP has a rank-one decomposition such that $\nabla_{W_\ell} L = \sum_{i=1}^B \delta_{\ell+1}^i u_\ell^{i\top}$ for a batch B , (Mitchell et al., 2021) are able to construct an editor network g_ℓ to generate the weight updates. Here, $\delta_{\ell+1}^i$ is the gradient for element i for the preactivations of layer $\ell + 1$ and u_ℓ^i are the inputs of element i into layer ℓ . To characterize these updates, MEND employs functions that map $\delta_{\ell+1}^i$ and u_ℓ^i to a pseudo-decomposition $\tilde{\delta}_{\ell+1}^i$ and \tilde{u}_ℓ^i such that $\tilde{\nabla}_{W_\ell} L = \sum_{i=1}^B \tilde{\delta}_{\ell+1}^i \tilde{u}_\ell^{i\top}$. Letting $z_\ell = \text{concat}(\delta_{\ell+1}, u_\ell)$, the form of the network is

$$h_\ell = z_\ell + \sigma(s_\ell^1 \odot (U_1 V_1 z_\ell + b) + o_\ell^1) \quad (9)$$

$$g(z_\ell) = h_\ell + \sigma(s_\ell^2 \odot U_2 V_2 h_\ell + o_\ell^2) \quad (10)$$

where σ is the ReLu activation function and U_j, V_j are a low rank decomposition of MEND’s weight for layer j . Note that, because of the difference in dimensions between weight matrices across layers, MEND learns different parameters for each unique shape of weight matrices to be edited. Additionally, layer-wise offset and scale parameters o_ℓ and s_ℓ are learned for both h_ℓ and g_ℓ . The final update is given by $\tilde{W} = W - \alpha_\ell \tilde{\nabla}_{W_\ell}$ with α_ℓ being another learned parameter per layer.

Given the original weights W and the updated weights \tilde{W} , loss is computed aggregating two training losses, editing success and locality. For a desired edit (x_e, y_e) , (x'_e, y'_e) is defined as a semantically equivalent wording of the edit. Editing loss is defined as $L_e = -\log_{p_{\theta_{\tilde{W}}}} p(y'_e | x'_e)$. x_{loc} is defined as a locality sample, which is randomly sampled to test the edited model’s impact on information

unrelated to the edit. The corresponding locality loss is $L_{\text{loc}} = D_{KL}(p_{\theta_W}(\cdot | x_{\text{loc}}) \| p_{\theta_{\tilde{W}}}(\cdot | x_{\text{loc}}))$.

The total loss is computed as $L_{\text{MEND}} = c_e L_e(\theta_{\tilde{W}}) + L_{\text{loc}}(\theta_W, \theta_{\tilde{W}})$ where $c_e = 0.1$. The total loss is optimized using the Adam optimizer. For GPT2-XL, we edit layers 46, 47, and 48. For GPTJ, we edit layers 26, 27, and 28.

A.6.4 FT

The Fine-Tuning procedures used in this paper follow from (Meng et al., 2022b) and (Meng et al., 2022a)’s implementation for both GPT-J and GPT2-XL. MLP weights for a single layer are fine-tuned for both models. We use a constrained fine tuning approach where we add a L_∞ constraint such that $\|\theta_G - \theta_{G'}\|_\infty \leq \epsilon$ at each gradient step. For the constraint, $\epsilon = 5e-4$ for GPT2-XL and $\epsilon = 5e-5$ for GPT-J. It is optimized using Adam with a learning rate of 5e-4 for both GPT2-XL and GPT-J. We fine tune layers 18 and 6 for GPT2-XL and GPT-J respectively.